

String matching

July 24, 2015

1 The scan left function

1.1 scan left

`scanl` is a polymorphic function of type $(\beta \rightarrow \alpha \rightarrow \beta) \rightarrow \beta \rightarrow [\alpha] \rightarrow \beta$ for type variables, α and β . It is defined by the following equations.

$$\begin{aligned} \text{scanl } f \ z \ [] &= [z] \\ \text{scanl } f \ z \ (x : xs) &= z : \text{scanl } f \ (f \ z \ x) \ xs \end{aligned}$$

In C++, lists are replaced with iterators.

```
template <
    class F, class AccT, class InT, class OutT>
OutT scan_left (F f, AccT z, InT begin, InT end, OutT out) {
    *out++ = z;
    if (begin == end) return out;
    auto const& x = *begin;

    return scan_left (f, f (z, x), std::next (begin), end, out);
}
```

2 The string matching problem

2.1 matches

A string matching problem is one in which one finds all occurrences of a non-empty string (the pattern) in some other string (the text). A specification for the problem can be stated like this:

$$\text{matches } ws = \text{map } \text{len} \cdot \text{filter } (\text{endswith } ws) \cdot \text{inits} \quad (1)$$

The function *inits* returns a list of the prefixes of the text in order of increasing length. The expression *endswith* *ws* *xs* tests whether the pattern *ws* is a suffix of *xs*. Finally, the value *matches* *ws* *xs* is a list of integers *p* such that *ws* is a suffix

of *take p xs*. For example: *matches "abcab" "ababcab"* is the list $[7, 10]$. That is, *matches ws xs* returns a list of integers p such that *ws* appears in *xs* ending at position p (counting positions from 1).

2.2 filter

In C++, *filter* can be written like this.

```
template <class PredT, class RngT, class OutT>
OutT filter (PredT p, RngT xs, OutT out) {
    return std::accumulate (
        std::begin (xs), std::end (xs), out,
        [&p](auto dst, auto const& x) {
            return p (x) ? *dst++ = x : dst;
        }
    );
}
```

2.3 endswith

We define *endswith ws* = $(reverse\ ws \sqsubseteq) \cdot reverse$ where \sqsubseteq is the *prefix* relation given by the equations

$$\begin{aligned} [] &\sqsubseteq us = true \\ (u : us) &\sqsubseteq [] = false \\ (u : us) &\sqsubseteq (v : vs) = (u = v \wedge us \sqsubseteq vs) \end{aligned}$$

Here's the *prefix* function realized in C++.

```
template <class InT1, class InT2>
bool prefix (InT1 lb, InT1 le, InT2 rb, InT2 re) {
    if (lb == le) return true;
    if (rb == re) return false;

    return *lb == *rb &&
        prefix (std::next (lb), le, std::next (rb), re);
}
```

3 The Boyer-Moore solution

3.1 Theory

This identity is called "the scan lemma".

$$map (foldl\ op\ e) \cdot inits = scanl\ op\ e \tag{2}$$

This equation is important because the left hand side has complexity $O(N^2)$ whereas the right-hand-side, $O(N)$. It admits restating equation 1 as:

$$\begin{aligned} matches\ ws &= map\ fst \cdot filter((sw \sqsubseteq) \cdot snd)\ scanl\ step\ (0, []) \\ sw &= reverse\ ws \\ step\ (n, sx)\ x &= (n + 1, x : sx) \end{aligned} \tag{3}$$

This is the called the basic “Boyer-Moore” algorithm. Here's an implmentation.

```
template <class OutT>
OutT matches (std::string const& ws, std::string const& s, OutT dst) {
    typedef std::pair<int, std::deque<char>> acc_t;

    auto step = [](acc_t p, char x) -> acc_t {
        ++p.first;
        p.second.push_front (x);

        return p;
    };

    std::deque<acc_t> buf;
    scan_left (
        step
        , std::make_pair(0, std::deque<char>())
        , s.begin ()
        , s.end ()
        , std::back_inserter(buf));

    std::string sw(ws.rbegin (), ws.rend ());
    auto pred = [&sw] (auto p) -> bool {
        return prefix (
            sw.begin (), sw.end ()
            , p.second.begin (), p.second.end ());
    };
    std::deque<acc_t> temp;
    filter (pred, buf, std::back_inserter (temp));

    return std::transform (
        temp.begin (), temp.end (), dst,
        [](acc_t const& p) -> int { return p.first; });
}
```

4 Testing

Try a basic test.

```

int main () {

    std::list<int> where;
    matches ("abcab", "ababcbcab", std::back_inserter (where));

    std::for_each (where.begin (), where.end ()
        , [](int i) -> void { std::cout << i << ", "; }
        );

    return 0;
}

```

This program should print "7, 10".

5 War and Peace

We can try the Boyer-Moore as implemented above in earnest by trying to find all occurrence of the string "people" in the text of the book "War and Peace" by Tolstoy. When we do, we'll immediately encounter the first problem - stack overflow. I'd expect the compiler to optimize the tail-call in `scan_left` but alas it does not. We can take matters into our own hands like so :

```

template <
    class F, class AccT, class InT, class OutT>
OutT scan_left (F f, AccT z, InT begin, InT end, OutT out) {
loop:
    *out++ = z;
    if (begin == end) return out;
    auto const& x = *begin;
    z = f (z, x);
    ++begin;
    goto loop;
}

```

What we've done so far is address stack overflow issues but we still have a problem with ridiculous memory consumption. This program will exhaust all available memory given this amount of input. `scan_left` is producing a list of `scan_left` of `deque<char>` of lengths $1 \dots N$ where N is the string length. There's no need to be allocating memory and copying this data. We can simply reference the character ranges in the source string using iterators.

```

template <class OutT>
OutT matches (std::string const& ws, std::string const& s, OutT dst) {
    typedef std::string::const_reverse_iterator it;
    typedef std::pair<it, it> iterator_range;
    typedef std::pair<int, iterator_range> acc_t;

```

```

std::size_t num_chars=s.size();

auto step = [num_chars,&s](acc_t p, char x) -> acc_t {
    ++p.first;
    std::string::const_reverse_iterator rbegin = s.rbegin ();
    std::advance (rbegin, num_chars - p.first);
    p.second = std::make_pair (rbegin, s.rend ());

    return p;
};

std::deque<acc_t> buf1;
scan_left (
    step
    , std::make_pair (0, std::make_pair (s.rend (), s.rend()))
    , s.begin ()
    , s.end ()
    , std::back_inserter (buf1));

std::string sw(ws.rbegin (), ws.rend ());
auto pred = [num_chars, &sw, &s] (auto p) -> bool {
    return prefix (sw.begin (), sw.end (), p.second.first, p.second.second);
};

std::deque<acc_t> buf2;
filter (pred, buf1, std::back_inserter (buf2));
buf2.swap (buf1);

return std::transform (buf1.begin (), buf1.end (),
    dst, [](acc_t const& p) -> int { return p.first; });
}

```

That's enough to solve the memory problems and now we can get an answer to the war and peace problem (582 occurrences by the way) but we can do better. As we've now realized that we can construct references to the character ranges via iterators, we realize that this can be done on demand. Therefore, we can represent the entire computation as a simple loop over the text eliminating `scan_left` and `filter` from the code entirely. Of course, they are still there but now only conceptually!

```

template <class OutT>
OutT matches2 (std::string const& ws, std::string const& s, OutT dst) {
    std::string sw (ws.rbegin (), ws.rend ());
    std::size_t num_chars=s.size ();
    for (std::size_t i = 0; i < num_chars; ++i) {
        std::string::const_reverse_iterator rbegin = s.rbegin ();
        std::advance (rbegin, num_chars - i);
        if (prefix (sw.begin (), sw.end (), rbegin, s.rend ())) {

```

```
        *dst++ = i;
    }
}

return dst;
}
```

References

- [1] Richard Bird, *Pearls of Functional Algorithm Design*. Cambridge University Press, 2010.